Effects of Logistic Regressions and Standard Scaler on Advertisements Data Set

By Aidan Ben-Josef

The problem

The goal is use the given data set to predict the likelihood of customers purchasing ads for the given features.

To analyze the data, a logistic regression was run on both the unscaled and scaled versions of the data.

Graphed Results

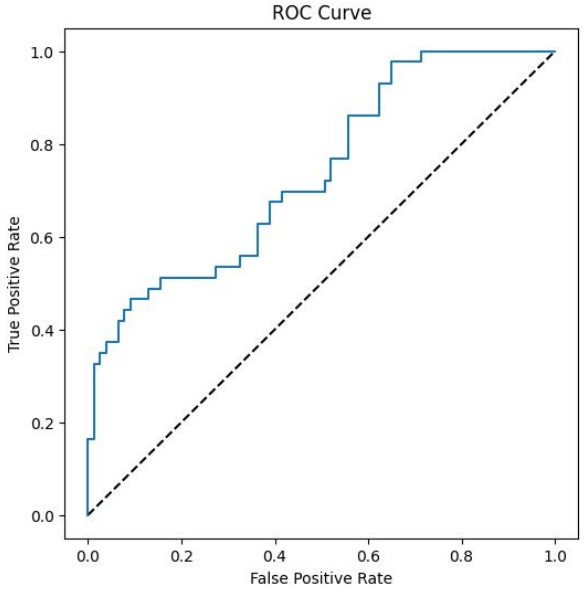

Unscaled Data

AUC = 0.7375415282392026

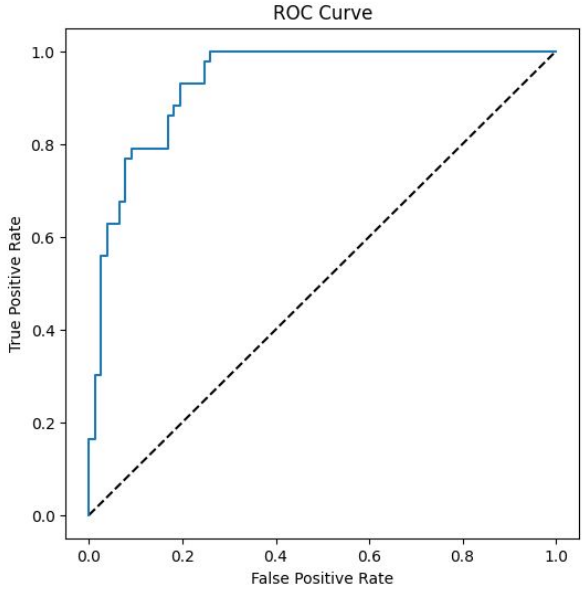

Scaled Data

AUC = 0.933856840833585

Conclusions

The unscaled data was useful, but could only be correct 73% of the time.

The scaled version is much more accurate, presenting a score that is correct 93% of the time.